

Lab 4 - Predicting Breast Cancer

Submitted By Name: Harsha KG Register Number: 19112005 Class: 5 BSc Data Science

Objectives

- Understand about Logistic Regression, K-Nearest Neighbours and Decision Trees

Problem Definition

Compare and Contrast the Differences in Classification Result among Logistic Regression, K-Nearest Neighbors and Decision Trees with respect to the Breast Cancer Dataset. Download the dataset as available in the URL: <https://www.kaggle.com/uciml/breast-cancer-wisconsin-data>
(<https://www.kaggle.com/uciml/breast-cancer-wisconsin-data>).

Demonstrate various evaluation metrics Check the effect of classification with respect to change in train-test dataset, classification parameters, hyper parameters etc

Use appropriate visualizations and interpretations on result cases Do the standard practices, as discussed in class

Approach

Imported the dataset using the required library. Did some preprocessing techniques before exploration to make the dataset into standard format and then did some EDA. After that dataset is splitted into training and testing set. Build models of Logistic, K Nearest Neighbour and Decision tree by changing the parameter values of each algorithm.

References:

1. Scikit Documentation
2. <https://towardsdatascience.com/machine-learning-with-python-classification-complete-tutorial-d2c99dc524ec> (<https://towardsdatascience.com/machine-learning-with-python-classification-complete-tutorial-d2c99dc524ec>).
3. <https://stackabuse.com/overview-of-classification-methods-in-python-with-scikit-learn/> (<https://stackabuse.com/overview-of-classification-methods-in-python-with-scikit-learn/>). [Evaluation Metrics]

In [1]:

```
# Importing the Libraries
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn import metrics
import hvplot.pandas
import itertools
import plotly.graph_objs as go
import plotly.tools as tls
import plotly.figure_factory as ff
import plotly.offline as py

from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split

import warnings
warnings.filterwarnings("ignore")
```

In [2]:

```
BC = pd.read_csv("Breast_Cancer_M14.csv")
```

In [3]:

```
BC.head()
```

Out[3]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothnes
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	

5 rows × 32 columns

In [4]:

```
BC.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 569 entries, 0 to 568
```

```
Data columns (total 32 columns):
```

#	Column	Non-Null Count	Dtype
0	id	569 non-null	int64
1	diagnosis	569 non-null	object
2	radius_mean	569 non-null	float64
3	texture_mean	569 non-null	float64
4	perimeter_mean	569 non-null	float64
5	area_mean	569 non-null	float64
6	smoothness_mean	569 non-null	float64
7	compactness_mean	569 non-null	float64
8	concavity_mean	569 non-null	float64
9	concave points_mean	569 non-null	float64
10	symmetry_mean	569 non-null	float64
11	fractal_dimension_mean	569 non-null	float64
12	radius_se	569 non-null	float64
13	texture_se	569 non-null	float64
14	perimeter_se	569 non-null	float64
15	area_se	569 non-null	float64
16	smoothness_se	569 non-null	float64
17	compactness_se	569 non-null	float64
18	concavity_se	569 non-null	float64
19	concave points_se	569 non-null	float64
20	symmetry_se	569 non-null	float64
21	fractal_dimension_se	569 non-null	float64
22	radius_worst	569 non-null	float64
23	texture_worst	569 non-null	float64
24	perimeter_worst	569 non-null	float64
25	area_worst	569 non-null	float64
26	smoothness_worst	569 non-null	float64
27	compactness_worst	569 non-null	float64
28	concavity_worst	569 non-null	float64
29	concave points_worst	569 non-null	float64
30	symmetry_worst	569 non-null	float64
31	fractal_dimension_worst	569 non-null	float64

```
dtypes: float64(30), int64(1), object(1)
```

```
memory usage: 142.4+ KB
```

In [5]:

```
BC.isnull().sum()  
BC.isna().sum()
```

Out[5]:

```
id                0  
diagnosis         0  
radius_mean       0  
texture_mean      0  
perimeter_mean    0  
area_mean         0  
smoothness_mean   0  
compactness_mean  0  
concavity_mean    0  
concave points_mean 0  
symmetry_mean     0  
fractal_dimension_mean 0  
radius_se         0  
texture_se        0  
perimeter_se      0  
area_se           0  
smoothness_se     0  
compactness_se    0  
concavity_se      0  
concave points_se 0  
symmetry_se       0  
fractal_dimension_se 0  
radius_worst      0  
texture_worst     0  
perimeter_worst   0  
area_worst        0  
smoothness_worst  0  
compactness_worst 0  
concavity_worst   0  
concave points_worst 0  
symmetry_worst    0  
fractal_dimension_worst 0  
dtype: int64
```

In [6]:

```
BC['diagnosis'].value_counts()
```

Out[6]:

```
B    357  
M    212  
Name: diagnosis, dtype: int64
```


In [7]:

```
#-----PERCENTAGE-----
trace = go.Pie(labels = ['benign', 'malignant'], values = BC['diagnosis'].value_counts(
),
               textfont=dict(size=15), opacity = 0.8,
               marker=dict(colors=['lightskyblue', 'gold'],
                           line=dict(color='#000000', width=1.5)))

layout = dict(title = 'Distribution of diagnosis variable')

fig = dict(data = [trace], layout=layout)
py.iplot(fig)
```

Distribution of diagnosis variable



In []:

In [8]:

```
#correlation
correlation = BC.corr()
#tick labels
matrix_cols = correlation.columns.tolist()
#convert to array
corr_array = np.array(correlation)
```

In [9]:

```
#Plotting
trace = go.Heatmap(z = corr_array,
                  x = matrix_cols,
                  y = matrix_cols,
                  xgap = 2,
                  ygap = 2,
                  colorscale='Viridis',
                  colorbar = dict() ,
                  )
layout = go.Layout(dict(title = 'Correlation Matrix for variables',
                        autosize = False,
                        height = 720,
                        width = 800,
                        margin = dict(r = 0 ,l = 210,
                                      t = 25,b = 210,
                                      ),
                        yaxis = dict(tickfont = dict(size = 9)),
                        xaxis = dict(tickfont = dict(size = 9)),
                        )
fig = go.Figure(data = [trace],layout = layout)
py.ipplot(fig)
```

Correlation Matrix for variables



In [10]:


```
plt.figure(figsize = (20,10))
sns.set_theme(style="darkgrid")

radius = BC[['radius_mean', 'radius_se', 'radius_worst', 'diagnosis']]
sns.pairplot(radius, hue='diagnosis', markers=["o", "s"])
```

Out[10]:

<seaborn.axisgrid.PairGrid at 0x1a08f50ec10>

<Figure size 1440x720 with 0 Axes>



In [11]:


```
plt.figure(figsize = (20,10))
sns.set_theme(style="darkgrid")

radius = BC[['radius_mean', 'radius_se', 'radius_worst', 'diagnosis']]
sns.pairplot(radius, hue='diagnosis', markers=["o", "s"])
```

Out[11]:

<seaborn.axisgrid.PairGrid at 0x1a0b7b36970>

<Figure size 1440x720 with 0 Axes>



In []:

In []:

In [14]:

```
# 2 datasets
M = BC[(BC['diagnosis'] != 0)]
B = BC[(BC['diagnosis'] == 0)]
```


In [15]:

```
def plot_feat1_feat2(feet1, feat2) :  
    trace0 = go.Scatter(  
        x = M[feat1],  
        y = M[feat2],  
        name = 'malignant',  
        mode = 'markers',  
        marker = dict(color = '#FFD700',  
            line = dict(  
                width = 1)))  
  
    trace1 = go.Scatter(  
        x = B[feat1],  
        y = B[feat2],  
        name = 'benign',  
        mode = 'markers',  
        marker = dict(color = '#7EC0EE',  
            line = dict(  
                width = 1)))  
  
    layout = dict(title = feat1 + " "+"vs"+" "+ feat2,  
        yaxis = dict(title = feat2,zeroline = False),  
        xaxis = dict(title = feat1, zeroline = False)  
    )  
  
    plots = [trace0, trace1]  
  
    fig = dict(data = plots, layout=layout)  
    py.iplot(fig)
```


In [16]:

```
plot_feat1_feat2('perimeter_mean', 'radius_worst')  
plot_feat1_feat2('area_mean', 'radius_worst')  
plot_feat1_feat2('texture_mean', 'texture_worst')  
plot_feat1_feat2('area_worst', 'radius_worst')
```


perimeter_mean vs radius_worst




area_mean vs radius_worst



texture_mean vs texture_worst




area_worst vs radius_worst




In [17]:

```
plot_feat1_feat2('smoothness_mean', 'texture_mean')  
plot_feat1_feat2('radius_mean', 'fractal_dimension_worst')  
plot_feat1_feat2('texture_mean', 'symmetry_mean')  
plot_feat1_feat2('texture_mean', 'symmetry_se')
```



smoothness_mean vs texture_mean




radius_mean vs fractal_dimension_worst



texture_mean vs symmetry_mean




texture_mean vs symmetry_se



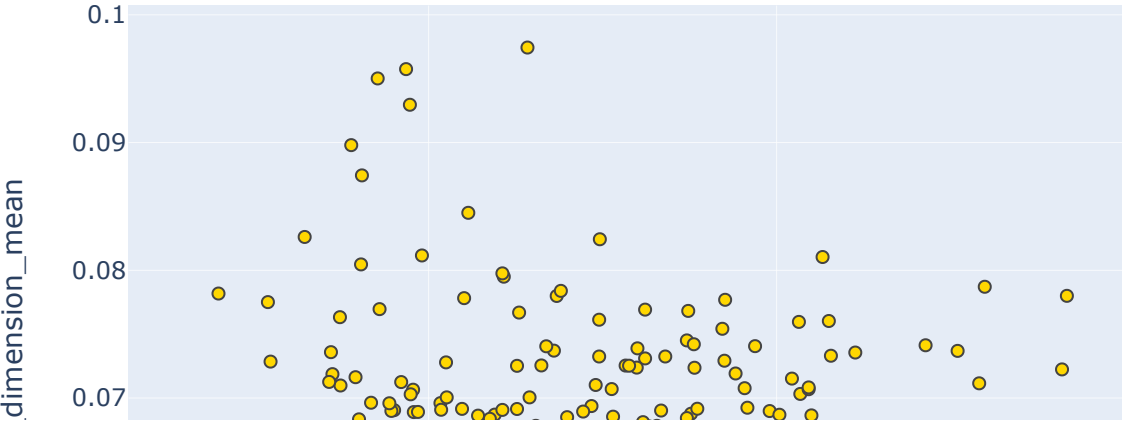
In [18]:

```
plot_feat1_feat2('area_mean', 'fractal_dimension_mean')  
plot_feat1_feat2('radius_mean', 'fractal_dimension_mean')  
plot_feat1_feat2('area_mean', 'smoothness_se')  
plot_feat1_feat2('smoothness_se', 'perimeter_mean')
```

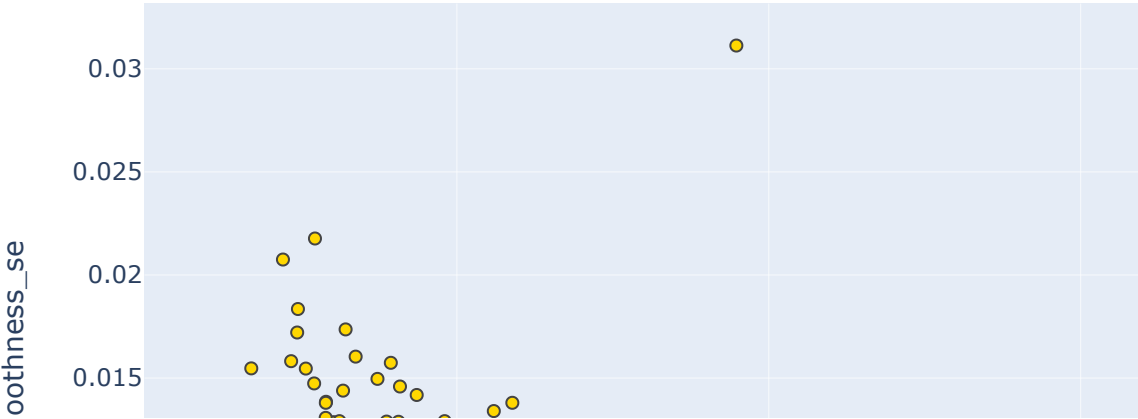
area_mean vs fractal_dimension_mean




radius_mean vs fractal_dimension_mean



area_mean vs smoothness_se




smoothness_se vs perimeter_mean



In [19]:

```
sns.barplot(x="radius_mean", y="texture_mean", data=BC[170:180])  
plt.title("Radius Mean vs Texture Mean", fontsize=15)  
plt.xlabel("Radius Mean")  
plt.ylabel("Texture Mean")  
plt.show()  
plt.style.use("ggplot")
```



In [20]:

```
BC.drop(columns='id',axis=1,inplace=True)
```

In [21]:

```
BC.head()
```

Out[21]:

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean
0	M	17.99	10.38	122.80	1001.0	0.11840	0.26340
1	M	20.57	17.77	132.90	1326.0	0.08474	0.26340
2	M	19.69	21.25	130.00	1203.0	0.10960	0.26340
3	M	11.42	20.38	77.58	386.1	0.14250	0.26340
4	M	20.29	14.34	135.10	1297.0	0.10030	0.26340

5 rows × 8 columns

In [22]:

```
Y = BC['diagnosis']  
X = BC.drop(['diagnosis'], axis=1)
```

In [23]:

```
Y.head()
```

Out[23]:

```
0    M  
1    M  
2    M  
3    M  
4    M  
Name: diagnosis, dtype: object
```

In [24]:

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = .30, random_state  
= 8)
```

In [25]:

```
X_train.head()
```

Out[25]:

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactne
547	10.26	16.58	65.85	320.8	0.08877	
331	12.98	19.35	84.52	514.0	0.09579	
421	14.69	13.98	98.22	656.1	0.10310	
95	20.26	23.03	132.40	1264.0	0.09078	
125	13.85	17.21	88.44	588.7	0.08785	

5 rows × 30 columns

In [26]:

```
X_test.head()
```

Out[26]:

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactne
325	12.670	17.30	81.25	489.9	0.10280	
557	9.423	27.88	59.26	271.3	0.08123	
475	12.830	15.73	82.89	506.9	0.09040	
308	13.500	12.71	85.69	566.2	0.07376	
553	9.333	21.94	59.01	264.0	0.09240	

5 rows × 30 columns

In [27]:

```
Y_test.head()
```

Out[27]:

```
325    B
557    B
475    B
308    B
553    B
Name: diagnosis, dtype: object
```

In [28]:

```
Y_train.head()
```

Out[28]:

```
547    B
331    B
421    B
95     M
125    B
Name: diagnosis, dtype: object
```

LogisticRegression

In [29]:

```
#Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

In [30]:

```
from sklearn.linear_model import LogisticRegression
```

In [31]:

```
logistic_regressor = LogisticRegression()
```

In [32]:

```
logistic_regressor.fit(X_train, Y_train)
```

Out[32]:

```
LogisticRegression()
```

In [33]:

```
y_pred = logistic_regressor.predict(X_test)
y_prob = logistic_regressor.predict_proba(X_test)
```

In [34]:

```
y_pred[0:5]
```

Out[34]:

```
array(['B', 'B', 'B', 'B', 'B'], dtype=object)
```

In [35]:

```
data = pd.DataFrame({'Actual': Y_test, 'Predicted': y_pred})  
data
```

Out[35]:

	Actual	Predicted
325	B	B
557	B	B
475	B	B
308	B	B
553	B	B
...
154	B	B
208	B	B
311	B	B
329	M	M
282	M	M

171 rows × 2 columns

In [36]:

```
y_prob[0:5]
```

Out[36]:


```
array([[9.98588993e-01, 1.41100703e-03],  
       [9.99647709e-01, 3.52291383e-04],  
       [9.92574976e-01, 7.42502412e-03],  
       [9.99968116e-01, 3.18836820e-05],  
       [9.99978480e-01, 2.15199463e-05]])
```

In [37]:

```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_test, y_pred)
dataframe_conf_matrix = conf_matrix
sns.heatmap(dataframe_conf_matrix, annot=True)
```

Out[37]:

<matplotlib.axes._subplots.AxesSubplot at 0x1a093de3070>



In [38]:

```
from sklearn.metrics import accuracy_score
acc_score1 = accuracy_score(Y_test, y_pred)
print(acc_score1)
```

0.9766081871345029

In [39]:

```
print("Training Score: ", logistic_regressor.score(X_train, Y_train)*100)
```

Training Score: 98.99497487437185

In [40]:

```
print("Testing Score: ", logistic_regressor.score(X_test, Y_test)*100)
```

Testing Score: 97.6608187134503

In [41]:

```
from sklearn.metrics import classification_report
class_report = classification_report(Y_test, y_pred)
print(class_report)
```

	precision	recall	f1-score	support
B	0.98	0.98	0.98	105
M	0.97	0.97	0.97	66
accuracy			0.98	171
macro avg	0.98	0.98	0.98	171
weighted avg	0.98	0.98	0.98	171

In [42]:

```
from sklearn.metrics import confusion_matrix
Cm = confusion_matrix(Y_test,y_pred)
Cm
```

Out[42]:

```
array([[103,  2],
       [ 2,  64]], dtype=int64)
```

In [43]:

```
Accuracy = (Cm[0][0] + Cm[1][1]) / (Cm[0][0] + Cm[1][1] + Cm[0][1] + Cm[1][0])
print("Accuracy",Accuracy)
Error_rate = (Cm[0][1] + Cm[1][0]) / (Cm[0][0] + Cm[1][1] + Cm[0][1] + Cm[1][0])
print("Error_rate",Error_rate)
Sensitivity = Cm[0][0]/(Cm[0][0] + Cm[1][0])
print("Sensitivity",Sensitivity)
Specificity = Cm[1][1]/(Cm[1][1] + Cm[0][1])
print("Specificity",Specificity)
Recall = Cm[0][0]/(Cm[0][0] + Cm[1][0])
print("Recall",Recall)
Precision = Cm[0][0]/(Cm[0][0] + Cm[0][1])
print("Precision",Precision)
F1Score = (2*(Precision*Recall))/(Precision + Recall)
print("F1Score",F1Score)
```

```
Accuracy 0.9766081871345029
Error_rate 0.023391812865497075
Sensitivity 0.9809523809523809
Specificity 0.9696969696969697
Recall 0.9809523809523809
Precision 0.9809523809523809
F1Score 0.9809523809523809
```

In [44]:

```
def doLogisticRegression(X, Y, test_size = 0.20, random_state = 42, penalty='l2', solver='lbfgs'):  
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = test_size, random_state = random_state)  
  
    logistic_regressor = LogisticRegression(penalty=penalty, solver=solver)  
    logistic_regressor.fit(X_train, Y_train)  
    y_pred = logistic_regressor.predict(X_test)  
  
    acc_score = accuracy_score(Y_test, y_pred)  
  
    return acc_score
```


In [45]:

```
penalties = ['none', 'l2']
test_size = [0.30, 0.25, 0.20]
random_states = [10, 25, 55]
solvers = ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga']

for t_size in test_size:
    for r_state in random_states:
        for penalty in penalties:
            for solver in solvers:
                accuracy = doLogisticRegression(X, Y, t_size, r_state, penalty)
                print("Test: {} | Random State: {} | Penalty: {} | Solver: {} | Accuracy : {}".format(t_size, r_state, penalty, solver, accuracy))
```

Test: 0.3 | Random State: 10 | Penalty: none | Solver: newton-cg | Accuracy : 0.9415204678362573
Test: 0.3 | Random State: 10 | Penalty: none | Solver: lbfgs | Accuracy : 0.9415204678362573
Test: 0.3 | Random State: 10 | Penalty: none | Solver: liblinear | Accuracy : 0.9415204678362573
Test: 0.3 | Random State: 10 | Penalty: none | Solver: sag | Accuracy : 0.9415204678362573
Test: 0.3 | Random State: 10 | Penalty: none | Solver: saga | Accuracy : 0.9415204678362573
Test: 0.3 | Random State: 10 | Penalty: 12 | Solver: newton-cg | Accuracy : 0.9473684210526315
Test: 0.3 | Random State: 10 | Penalty: 12 | Solver: lbfgs | Accuracy : 0.9473684210526315
Test: 0.3 | Random State: 10 | Penalty: 12 | Solver: liblinear | Accuracy : 0.9473684210526315
Test: 0.3 | Random State: 10 | Penalty: 12 | Solver: sag | Accuracy : 0.9473684210526315
Test: 0.3 | Random State: 10 | Penalty: 12 | Solver: saga | Accuracy : 0.9473684210526315
Test: 0.3 | Random State: 25 | Penalty: none | Solver: newton-cg | Accuracy : 0.9298245614035088
Test: 0.3 | Random State: 25 | Penalty: none | Solver: lbfgs | Accuracy : 0.9298245614035088
Test: 0.3 | Random State: 25 | Penalty: none | Solver: liblinear | Accuracy : 0.9298245614035088
Test: 0.3 | Random State: 25 | Penalty: none | Solver: sag | Accuracy : 0.9298245614035088
Test: 0.3 | Random State: 25 | Penalty: none | Solver: saga | Accuracy : 0.9298245614035088
Test: 0.3 | Random State: 25 | Penalty: 12 | Solver: newton-cg | Accuracy : 0.9239766081871345
Test: 0.3 | Random State: 25 | Penalty: 12 | Solver: lbfgs | Accuracy : 0.9239766081871345
Test: 0.3 | Random State: 25 | Penalty: 12 | Solver: liblinear | Accuracy : 0.9239766081871345
Test: 0.3 | Random State: 25 | Penalty: 12 | Solver: sag | Accuracy : 0.9239766081871345
Test: 0.3 | Random State: 25 | Penalty: 12 | Solver: saga | Accuracy : 0.9239766081871345
Test: 0.3 | Random State: 55 | Penalty: none | Solver: newton-cg | Accuracy : 0.9532163742690059
Test: 0.3 | Random State: 55 | Penalty: none | Solver: lbfgs | Accuracy : 0.9532163742690059
Test: 0.3 | Random State: 55 | Penalty: none | Solver: liblinear | Accuracy : 0.9532163742690059
Test: 0.3 | Random State: 55 | Penalty: none | Solver: sag | Accuracy : 0.9532163742690059
Test: 0.3 | Random State: 55 | Penalty: none | Solver: saga | Accuracy : 0.9532163742690059
Test: 0.3 | Random State: 55 | Penalty: 12 | Solver: newton-cg | Accuracy : 0.9532163742690059
Test: 0.3 | Random State: 55 | Penalty: 12 | Solver: lbfgs | Accuracy : 0.9532163742690059
Test: 0.3 | Random State: 55 | Penalty: 12 | Solver: liblinear | Accuracy : 0.9532163742690059
Test: 0.3 | Random State: 55 | Penalty: 12 | Solver: sag | Accuracy : 0.9532163742690059
Test: 0.3 | Random State: 55 | Penalty: 12 | Solver: saga | Accuracy : 0.9532163742690059
Test: 0.25 | Random State: 10 | Penalty: none | Solver: newton-cg | Accuracy :

cy : 0.9300699300699301
Test: 0.25 | Random State: 10 | Penalty: none | Solver: lbfgs | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 10 | Penalty: none | Solver: liblinear | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 10 | Penalty: none | Solver: sag | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 10 | Penalty: none | Solver: saga | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 10 | Penalty: 12 | Solver: newton-cg | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 10 | Penalty: 12 | Solver: lbfgs | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 10 | Penalty: 12 | Solver: liblinear | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 10 | Penalty: 12 | Solver: sag | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 10 | Penalty: 12 | Solver: saga | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 25 | Penalty: none | Solver: newton-cg | Accuracy : 0.916083916083916
Test: 0.25 | Random State: 25 | Penalty: none | Solver: lbfgs | Accuracy : 0.916083916083916
Test: 0.25 | Random State: 25 | Penalty: none | Solver: liblinear | Accuracy : 0.916083916083916
Test: 0.25 | Random State: 25 | Penalty: none | Solver: sag | Accuracy : 0.916083916083916
Test: 0.25 | Random State: 25 | Penalty: none | Solver: saga | Accuracy : 0.916083916083916
Test: 0.25 | Random State: 25 | Penalty: 12 | Solver: newton-cg | Accuracy : 0.916083916083916
Test: 0.25 | Random State: 25 | Penalty: 12 | Solver: lbfgs | Accuracy : 0.916083916083916
Test: 0.25 | Random State: 25 | Penalty: 12 | Solver: liblinear | Accuracy : 0.916083916083916
Test: 0.25 | Random State: 25 | Penalty: 12 | Solver: sag | Accuracy : 0.916083916083916
Test: 0.25 | Random State: 25 | Penalty: 12 | Solver: saga | Accuracy : 0.916083916083916
Test: 0.25 | Random State: 55 | Penalty: none | Solver: newton-cg | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 55 | Penalty: none | Solver: lbfgs | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 55 | Penalty: none | Solver: liblinear | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 55 | Penalty: none | Solver: sag | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 55 | Penalty: none | Solver: saga | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 55 | Penalty: 12 | Solver: newton-cg | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 55 | Penalty: 12 | Solver: lbfgs | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 55 | Penalty: 12 | Solver: liblinear | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 55 | Penalty: 12 | Solver: sag | Accuracy : 0.9300699300699301
Test: 0.25 | Random State: 55 | Penalty: 12 | Solver: saga | Accuracy : 0.9300699300699301
Test: 0.2 | Random State: 10 | Penalty: none | Solver: newton-cg | Accuracy : 0.9385964912280702

Test: 0.2 | Random State: 10 | Penalty: none | Solver: lbfgs | Accuracy : 0.9385964912280702
Test: 0.2 | Random State: 10 | Penalty: none | Solver: liblinear | Accuracy : 0.9385964912280702
Test: 0.2 | Random State: 10 | Penalty: none | Solver: sag | Accuracy : 0.9385964912280702
Test: 0.2 | Random State: 10 | Penalty: none | Solver: saga | Accuracy : 0.9385964912280702
Test: 0.2 | Random State: 10 | Penalty: 12 | Solver: newton-cg | Accuracy : 0.9298245614035088
Test: 0.2 | Random State: 10 | Penalty: 12 | Solver: lbfgs | Accuracy : 0.9298245614035088
Test: 0.2 | Random State: 10 | Penalty: 12 | Solver: liblinear | Accuracy : 0.9298245614035088
Test: 0.2 | Random State: 10 | Penalty: 12 | Solver: sag | Accuracy : 0.9298245614035088
Test: 0.2 | Random State: 10 | Penalty: 12 | Solver: saga | Accuracy : 0.9298245614035088
Test: 0.2 | Random State: 25 | Penalty: none | Solver: newton-cg | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 25 | Penalty: none | Solver: lbfgs | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 25 | Penalty: none | Solver: liblinear | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 25 | Penalty: none | Solver: sag | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 25 | Penalty: none | Solver: saga | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 25 | Penalty: 12 | Solver: newton-cg | Accuracy : 0.9385964912280702
Test: 0.2 | Random State: 25 | Penalty: 12 | Solver: lbfgs | Accuracy : 0.9385964912280702
Test: 0.2 | Random State: 25 | Penalty: 12 | Solver: liblinear | Accuracy : 0.9385964912280702
Test: 0.2 | Random State: 25 | Penalty: 12 | Solver: sag | Accuracy : 0.9385964912280702
Test: 0.2 | Random State: 25 | Penalty: 12 | Solver: saga | Accuracy : 0.9385964912280702
Test: 0.2 | Random State: 55 | Penalty: none | Solver: newton-cg | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 55 | Penalty: none | Solver: lbfgs | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 55 | Penalty: none | Solver: liblinear | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 55 | Penalty: none | Solver: sag | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 55 | Penalty: none | Solver: saga | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 55 | Penalty: 12 | Solver: newton-cg | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 55 | Penalty: 12 | Solver: lbfgs | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 55 | Penalty: 12 | Solver: liblinear | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 55 | Penalty: 12 | Solver: sag | Accuracy : 0.9210526315789473
Test: 0.2 | Random State: 55 | Penalty: 12 | Solver: saga | Accuracy : 0.9210526315789473

In [46]:

```
BC_1= pd.DataFrame(columns = ['Test Size', 'Random States', 'Penalty', 'Solvers', 'Accuracy'])
BC_1
```

Out[46]:

Test Size	Random States	Penalty	Solvers	Accuracy
-----------	---------------	---------	---------	----------

In [48]:

```
penalties = ['none', 'l2']
test_size = [0.30, 0.25, 0.20]
random_states = [10, 25, 55]
solvers = ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga']
for t_size in test_size:
    for r_state in random_states:
        for penalty in penalties:
            for solver in solvers:
                accuracy = doLogisticRegression(X, Y, t_size, r_state, penalty)

    #print("Test: {} | Random State: {} | Penalty: {} | Solver: {} | Accuracy : {}".format
    (t_size, r_state, penalty, solver, accuracy))

    BCEvaluation = {}
    BCEvaluation['Test Size'] = t_size
    BCEvaluation['Random States'] = r_state
    BCEvaluation['Penalty'] = penalty
    BCEvaluation['Solvers'] = solver
    BCEvaluation['Accuracy'] = accuracy
    BC_1= BC_1.append(BCEvaluation, ignore_index = True)
```

In [49]:

```
BC_1
```

Out[49]:

	Test Size	Random States	Penalty	Solvers	Accuracy
0	0.3	10	none	newton-cg	0.941520
1	0.3	10	none	lbfgs	0.941520
2	0.3	10	none	liblinear	0.941520
3	0.3	10	none	sag	0.941520
4	0.3	10	none	saga	0.941520
...
175	0.2	55	l2	newton-cg	0.921053
176	0.2	55	l2	lbfgs	0.921053
177	0.2	55	l2	liblinear	0.921053
178	0.2	55	l2	sag	0.921053
179	0.2	55	l2	saga	0.921053


180 rows × 5 columns

In [50]:

```
sns.displot(x = 'Accuracy', data = BC_1)
```

Out[50]:

<seaborn.axisgrid.FacetGrid at 0x1a093ed4c10>



K-Nearest Neighbors

In [80]:

```
from sklearn.neighbors import KNeighborsClassifier
classifier= KNeighborsClassifier()
classifier.fit(X_train, Y_train)
```

Out[80]:

KNeighborsClassifier()

In [81]:

```
y_pred= classifier.predict(X_test)
```

In [82]:

```
from sklearn.metrics import accuracy_score
acc_score2 = accuracy_score(Y_test, y_pred)
print(acc_score2)
```


0.9590643274853801

In [83]:

```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_test, y_pred)
dataframe_conf_matrix = conf_matrix
sns.heatmap(dataframe_conf_matrix, annot=True)
```

Out[83]:

<matplotlib.axes._subplots.AxesSubplot at 0x1a0939d98b0>



In [84]:

```
Accuracy = (Cm[0][0] + Cm[1][1]) / (Cm[0][0] + Cm[1][1] + Cm[0][1] + Cm[1][0])
print("Accuracy",Accuracy)
Error_rate = (Cm[0][1] + Cm[1][0]) / (Cm[0][0] + Cm[1][1] + Cm[0][1] + Cm[1][0])
print("Error_rate",Error_rate)
Sensitivity = Cm[0][0]/(Cm[0][0] + Cm[1][0])
print("Sensitivity",Sensitivity)
Specificity = Cm[1][1]/(Cm[1][1] + Cm[0][1])
print("Specificity",Specificity)
Recall = Cm[0][0]/(Cm[0][0] + Cm[1][0])
print("Recall",Recall)
Precision = Cm[0][0]/(Cm[0][0] + Cm[0][1])
print("Precision",Precision)
F1Score = (2*(Precision*Recall))/(Precision + Recall)
print("F1Score",F1Score)
```

```
Accuracy 0.9766081871345029
Error_rate 0.023391812865497075
Sensitivity 0.9809523809523809
Specificity 0.9696969696969697
Recall 0.9809523809523809
Precision 0.9809523809523809
F1Score 0.9809523809523809
```


In [85]:

```
error_rate = []  
  
# Will take some time  
for i in range(1, 40):  
  
    knn = KNeighborsClassifier(n_neighbors = i)  
    knn.fit(X_train, Y_train)  
    pred_i = knn.predict(X_test)  
    error_rate.append(np.mean(pred_i != Y_test))  
  
plt.figure(figsize=(10, 6))  
plt.plot(range(1, 40), error_rate, color='blue',  
         linestyle='dashed', marker='o',  
         markerfacecolor='red', markersize=10)  
  
plt.title('Error Rate vs. K Value')  
plt.xlabel('K')  
plt.ylabel('Error Rate')
```

Out[85]:

Text(0, 0.5, 'Error Rate')



In [128]:

```
BC_2=pd.DataFrame(columns=['n_neighbour','leaf_size','Accuracy'])
```

In [129]:

```
BC_2
```

Out[129]:

n_neighbour	leaf_size	Accuracy
-------------	-----------	----------

In [154]:

```
def doKNeighborsClassifier(X, Y, test_size = 0.20, random_state = 42):
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = test_size, random_state = random_state)

    KNeighbors_Classifier = KNeighborsClassifier()
    KNeighbors_Classifier.fit(X_train, Y_train)
    y_pred = KNeighbors_Classifier.predict(X_test)

    acc_score = accuracy_score(Y_test, y_pred)
```

In [156]:

```
n_neighbour = [5, 25, 35]
leaf_size = [10, 25, 55]

for neighbour in n_neighbour:
    for leaf in leaf_size:
        accuracy = doKNeighborsClassifier(X, Y,neighbour,leaf)

        BCEvaluation_KNN = {}
        BCEvaluation_KNN['n_neighbour'] = neighbour
        BCEvaluation_KNN['leaf_size'] = leaf
        BCEvaluation_KNN['Accuracy'] = accuracy
        BC_2= BC_2.append(BCEvaluation_KNN, ignore_index = True)
```

Decision Tree

In [98]:

```
#Fitting Decision Tree classifier to the training set
from sklearn.tree import DecisionTreeClassifier
classifier= DecisionTreeClassifier()
classifier.fit(X_train, Y_train)
```

Out[98]:

```
DecisionTreeClassifier()
```

In [99]:

```
#Predicting the test set result
y_pred= classifier.predict(X_test)
y_pred
```

Out[99]:


```
array(['B', 'B', 'B', 'B', 'B', 'B', 'B', 'B', 'B', 'B', 'B', 'B', 'B', 'B',
      'M', 'M', 'B', 'B', 'M', 'B', 'M', 'M', 'M', 'M', 'B', 'B', 'M',
      'M', 'M', 'B', 'B', 'B', 'B', 'M', 'M', 'M', 'B', 'M', 'M', 'B',
      'M', 'B', 'M', 'M', 'B', 'M', 'B', 'M', 'M', 'M', 'M', 'B', 'M',
      'B', 'M', 'B', 'M', 'M', 'B', 'B', 'B', 'M', 'M', 'B', 'B', 'B',
      'B', 'B', 'M', 'B', 'M', 'M', 'B', 'M', 'B', 'M', 'B', 'M', 'B',
      'B', 'B', 'M', 'M', 'B', 'M', 'B', 'B', 'M', 'M', 'B', 'M', 'B',
      'B', 'B', 'B', 'B', 'M', 'B', 'B', 'B', 'M', 'B', 'M', 'B', 'M',
      'B', 'M', 'B', 'M', 'B', 'B', 'B', 'B', 'B', 'M', 'B', 'B', 'B',
      'B', 'M', 'B', 'M', 'B', 'M', 'B', 'B', 'B', 'B', 'B', 'M', 'M',
      'B', 'B', 'B', 'B', 'B', 'B', 'B', 'M', 'B', 'B', 'B', 'B', 'B',
      'M', 'M', 'M', 'B', 'M', 'B', 'M', 'M', 'M', 'B', 'B', 'B', 'B',
      'M', 'M'], dtype=object)
```

In [135]:

```
from sklearn.metrics import confusion_matrix
conf_matrix = confusion_matrix(Y_test, y_pred)
dataframe_conf_matrix = conf_matrix
sns.heatmap(dataframe_conf_matrix, annot=True)
```

Out[135]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x1a0942b3b20>
```



In [100]:

```
#Creating the Confusion matrix
from sklearn.metrics import confusion_matrix
cm= confusion_matrix(Y_test, y_pred)
cm
```

Out[100]:

```
array([[98,  7],
       [ 5, 61]], dtype=int64)
```

In [101]:

```
from sklearn.metrics import accuracy_score
acc_score3 = accuracy_score(Y_test, y_pred)
print(acc_score3)
```

0.9298245614035088

In [139]:

```
def doDecisionTreeClassifier(X, Y, test_size = 0.20, random_state = 42):
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = test_size, random_state = random_state)
    classifier= DecisionTreeClassifier()
    classifier.fit(X_train, Y_train)
    y_pred= classifier.predict(X_test)

    acc_score = accuracy_score(Y_test, y_pred)
```

In [143]:

```
BC_3= pd.DataFrame(columns = ['Test Size', 'Random States'])
BC_3
```

Out[143]:

Test Size	Random States
-----------	---------------

In [144]:

```
test_size = [0.30, 0.25, 0.20]
random_states = [10, 25, 55]
for t_size in test_size:
    for r_state in random_states:
        accuracy = doDecisionTreeClassifier(X, Y, t_size, r_state)

    #print("Test: {} | Random State: {} | Penalty: {} | Solver: {} | Accuracy : {}".format(
    t_size, r_state, penalty, solver, accuracy))

    BCEvaluation = {}
    BCEvaluation['Test Size'] = t_size
    BCEvaluation['Random States'] = r_state
    BCEvaluation['Accuracy'] = accuracy
    BC_3= BC_3.append(BCEvaluation, ignore_index = True)
```

Accuracy Score of Models

In [102]:

```
prediction_columns = ["NAME OF MODEL", "ACCURACY SCORE"]
df_pred = {"NAME OF MODEL" : ["LOGISTIC REGRESSION", "K-NN", "DECISION TREE"],
           "ACCURACY SCORE " : [acc_score1, acc_score2, acc_score3]}
df_predictions = pd.DataFrame (df_pred)
df_predictions
```

Out[102]:

	NAME OF MODEL	ACCURACY SCORE
0	LOGISTIC REGRESSION	0.976608
1	K-NN	0.959064
2	DECISION TREE	0.929825

Hyperparameter Tuning

In [138]:

```
parameters = [{'penalty': ['l1', 'l2'], 'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000],
               'solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga']}]
grid_search = GridSearchCV(estimator = logistic_regressor,
                           param_grid = parameters,
                           scoring = 'accuracy',
                           cv = 10,
                           n_jobs = -1)
grid_search.fit(X_train, Y_train)
best_accuracy_log = grid_search.best_score_
best_parameters = grid_search.best_params_
print(best_accuracy_log)
print(best_parameters)
```

0.9824999999999999

{ 'C': 0.1, 'penalty': 'l2', 'solver': 'liblinear' }

In [106]:

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
oHe = OneHotEncoder()
```

In [120]:

```

parameters = [{'criterion': ['gini', 'entropy'], 'max_depth': [4, 5, 6, 7, 8, 9, 10, 11, 12, 15, 20, 30, 40, 50, 70, 90, 120, 150],
                'max_leaf_nodes': [2, 4, 6, 10, 15, 30, 40, 50, 100], 'min_samples_split': [2, 3, 4]}]
ohe = OneHotEncoder()
grid_search = GridSearchCV(estimator = classifier,
                           param_grid = parameters,
                           scoring = 'accuracy',
                           cv = 10,
                           n_jobs = -1)
grid_search.fit(X_train, Y_train)
best_accuracy_dtc = grid_search.best_score_
best_parameters = grid_search.best_params_
print(best_accuracy_dtc)
print(best_parameters)

```

0.9246153846153847

```

{'criterion': 'gini', 'max_depth': 7, 'max_leaf_nodes': 30, 'min_samples_split': 4}

```

Comparing models before and after Parameter Tuning

In [116]:

```

prediction_columns = ["NAME OF MODEL", "ACCURACY SCORE", "BEST ACCURACY (AFTER HYPER-PARAMETER TUNING)"]
df_pred = {"NAME OF MODEL" : ["LOGISTIC REGRESSION", "K-NN", "DECISION TREE", ],
           "ACCURACY SCORE " : [acc_score1, acc_score2, acc_score3],
           "BEST ACCURACY (AFTER HYPER-PARAMETER TUNING)" : [best_accuracy_log, best_accuracy_knn, best_accuracy_dtc]}
df_predictions = pd.DataFrame (df_pred)
df_predictions

```

Out[116]:

	NAME OF MODEL	ACCURACY SCORE	BEST ACCURACY (AFTER HYPER-PARAMETER TUNING)
0	LOGISTIC REGRESSION	0.976608	0.982500
1	K-NN	0.959064	0.969872
2	DECISION TREE	0.929825	0.926987

Conclusion

In []:

To conclude this notebook, it **is** fairly evident that it **is** the Logistic Regression model that has come out triumphant **with** the highest accuracies both before **and** after the hyper-parameter tuning . It ended up **with** an accuracy of 97% before hyper-parameter tuning **and** 98% after **and is** hence, the best suited model out of the rest **for** the given dataset.